

Tiered data disclosure model

What each audience receives from one complaint record. Every view is produced server-side by a fixed projection; clients never receive the full document.

Attribute of a complaint	Administrator	Supervisor and field worker	Public registry no login	Public tracker by tracking ID	Researcher anonymised records	Researcher aggregate only
IDENTITY						
Citizen name, email, phone	●	×	×	×	×	×
				shown as “Citizen”		
Assigned and assigning staff	●	●	×	●	×	×
				names in timeline		
Record identifier	●	●	●	●	●	×
	tracking ID	tracking ID	tracking ID	tracking ID	HMAC-SHA256, 12 hex	
CONTENT						
Title and free-text description	●	●	●	●	×	×
Photos: citizen, AI-annotated, proof	●	●	●	●	×	×
Street-level location	●	●	●	●	●	×
					coarsened to ward	
GOVERNANCE ATTRIBUTES						
Category, urgency, status, ward	●	●	●	●	●	●
						grouped counts
Filing and closing time	●	●	●	●	●	●
					day precision	monthly series
Resolution time	●	●	●	●	●	●
					derived, days	averages
Status history with remarks	●	●	●	●	×	×
			without actors	citizen redacted		

● full value ● transformed (coarsened, pseudonymised or aggregated) × withheld

Least privilege by audience
Staff see what the work needs; triage hides the filer's identity because the issue, not the person, is being judged.

Public by default, personal never
Every complaint is published on filing with its full timeline; identity fields are removed by the query itself, not by the client.

Research data is a product, not a dump
Researchers receive a governed, logged, expiring projection that cannot be joined back to a citizen without the server secret.

RESEARCH ACCESS CONTROL CHAIN

